

# Pregel In Graphs

## Models & Instances

Chase Zhang

Strikingly

March 16, 2018

# Table of Contents

Computing Model

Examples

System Design

Instance: User Tracking System

# Computing Model

## What is Pregel?

- ▶ Published by Google at 2009<sup>1</sup>
- ▶ Distributed graph computing system at large scale
- ▶ Implemented by open source solutions like Spark's GraphX<sup>2</sup>

---

<sup>1</sup>[https://kowshik.github.io/JPregel/pregel\\_paper.pdf](https://kowshik.github.io/JPregel/pregel_paper.pdf)

<sup>2</sup><https://spark.apache.org/graphx/>

# Computing Model

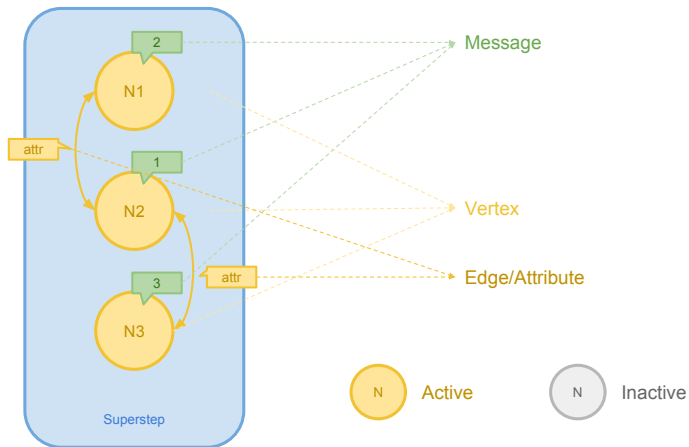



Figure: Basic Model

# Computing Model



**Figure:** Pregel is vertex oriented

# Computing Model

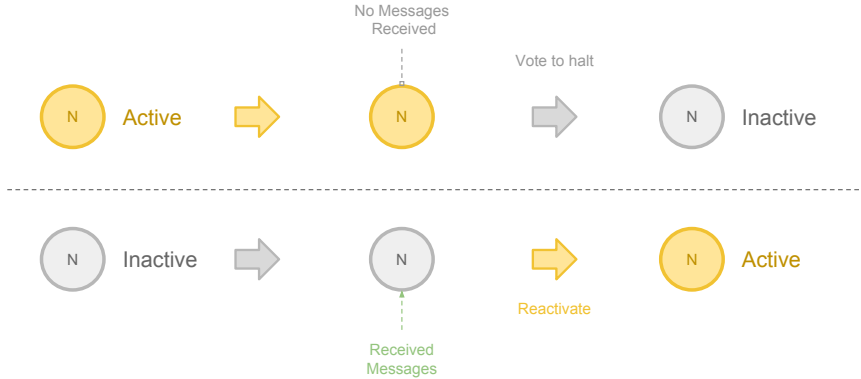


Figure: Vertices have states

# Computing Model

- ▶ Pregel is a vertex oriented computing model.
  - ▶ Map runs upon each vertex (other than edges)
  - ▶ Each vertex only knows its out-going edges during each superstep
- ▶ Pregel depends on message sending to iteratively computing the result
  - ▶ Each vertex receives messages from prior step
  - ▶ After one step, vertices send messages to adjacent vertices
- ▶ Vertices have states
  - ▶ Only active vertices will send message to the others
  - ▶ Program terminates once there is no active vertices

# Computing Model

## Problem

Sending message through network is costly.



# Computing Model




Figure: Passing messages

# Computing Model




Figure: Solution: Combiner

# Computing Model

## Problem

Some iterative graph algorithms require graph-wide results and metrics.

# Computing Model




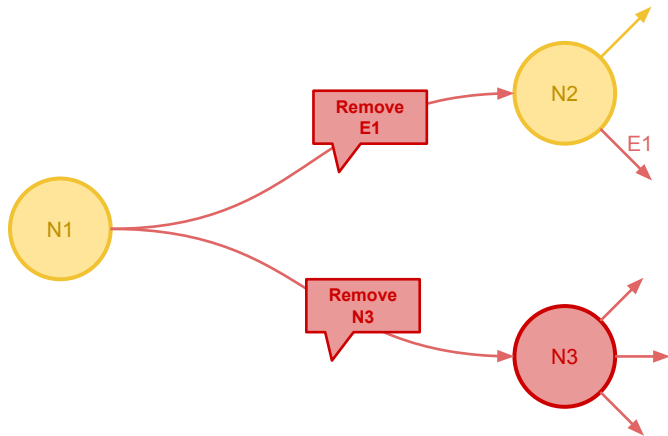
Figure: Solution: Aggregator

# Computing Model

## Problem

Some graph algorithms require modifying topology during iteration.

# Computing Model



**Figure:** Solution: Sending modification messages

# Computing Model

- ▶ Pregel provides **combiner** for reducing network traffic
- ▶ Pregel provides **aggregator** for recording graph-wide values and metrics
- ▶ Vertices in Pregel can send topology modification requests as messages to target vertices

# Examples

## Problem

Find connected components of a directed graph.



# Examples

## Solution

(Single machine) Union-Find<sup>3</sup>.

---

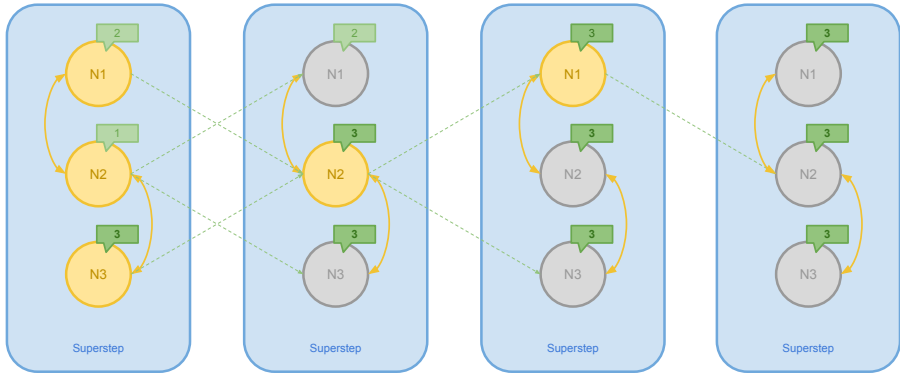
<sup>3</sup>[https://en.wikipedia.org/wiki/Disjoint-set\\_data\\_structure](https://en.wikipedia.org/wiki/Disjoint-set_data_structure)

# Examples

## Solution

(Pregel) Emit **max**(min) message received until no active vertex.

# Examples



**Figure:** Connected Components in Pregel

# Examples

## Problem

Find shortest path from a single source to the other vertices.

# Examples

## Solution

(Single machine) Dijkstra<sup>4</sup>, Bellman-Ford<sup>5</sup>, A\*<sup>6</sup>.

---

<sup>4</sup>[https://en.wikipedia.org/wiki/Dijkstra%27s\\_algorithm](https://en.wikipedia.org/wiki/Dijkstra%27s_algorithm)

<sup>5</sup>[https://en.wikipedia.org/wiki/Bellman%E2%80%93Ford\\_algorithm](https://en.wikipedia.org/wiki/Bellman%E2%80%93Ford_algorithm)

<sup>6</sup>[https://en.wikipedia.org/wiki/A\\*\\_search\\_algorithm](https://en.wikipedia.org/wiki/A*_search_algorithm)

# Examples

## Solution

(Pregel) BFS, emitting current shortest cost from source vertex to current vertex plus edge cost.

# Examples




Figure: Shortest Path: Step 1

# Examples




Figure: Shortest Path: Step 2



# Examples




Figure: Shortest Path: Step 3

# Examples

## Problem

PageRank<sup>7</sup>

---

<sup>7</sup><https://en.wikipedia.org/wiki/PageRank>

# Examples

$|V|$  vertices number

$|E(v)|$  edges number of vertex  $v$

$val(v)$  current message value of  $v$

$sum(v)$  sum of messages received of  $v$  in last superstep

- ▶ Initial messages:  $\frac{1}{|V|}$
- ▶ Emit messages:  $\frac{val(v)}{|E(v)|}$
- ▶ Update value:  $sum(v) \cdot 0.85 + \frac{0.15}{|V|}$

# Examples




Figure: PageRank

# Examples





Figure: PageRank

# System Design

- ▶ Graph partitioning
- ▶ Master-worker system model
- ▶ Message buffer for better performance
- ▶ Checkpoint and confined recovery

# System Design



**Figure:** Graph Partitioning

# System Design

- ▶ Graph is partitioned by `hash(vertexId)` by default
- ▶ User partitioning function can be provided for locality



# System Design




Figure: Master-Worker Model

# System Design

- ▶ Master coordinates supersteps, it holds no partition of graph
- ▶ Master calculate and hold values of aggregators
- ▶ Master send RPC to every participated worker and wait for task to finish
- ▶ Workers hold partitions of graph. Given a vertexId, a worker can know which worker is holding it without querying master

# System Design

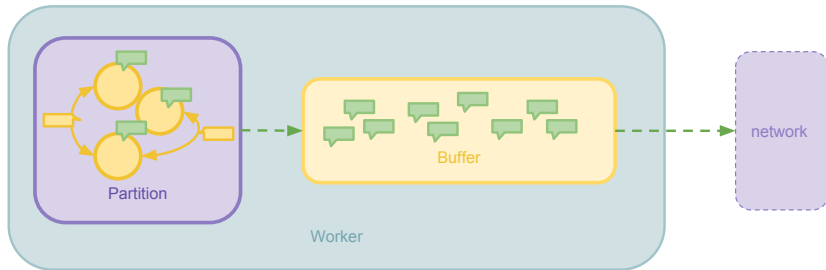



Figure: Message Buffer

# System Design



**Figure:** Checkpoint & Confined Recovery

# System Design

- ▶ A worker buffers message locally and send them as a batch as to reduce network overhead
- ▶ Before and after a superstep, a worker checkpoint and perform confined recovery
- ▶ Confined recovery logs every outgoing messages, so that only lost partitions need to be recalculated

# Instance: User Tracking System

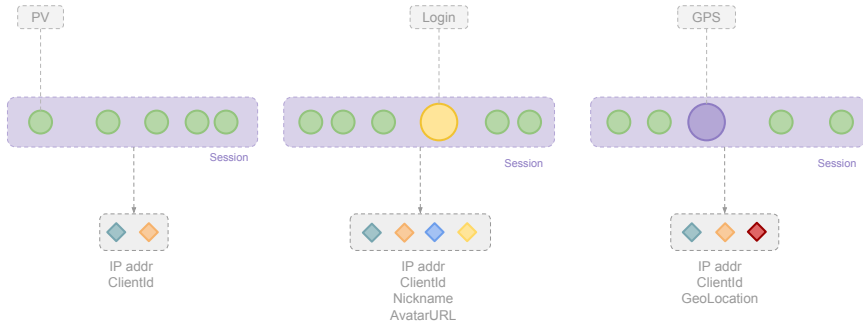




Figure: User Tracking Problem

# Instance: User Tracking System



**Figure:** Sessionize Events & Feature Extraction

# Instance: User Tracking System



**Figure:** Find Connected Components & Get User Profile



# Instance: User Tracking System

## Problem

Calculating edges between sessions cost  $O(N^2)$  time if we compare each pair. It will be too enormous even for just a million( $10^6$ ) sessions.

# Instance: User Tracking System

## Solution

Five significant features are selected for matching, and we define two sessions are matched if more than two features matched. So we can:

- ▶ Enumerate combinations of the five significant features, it will be  $C_5^2 = 10$
- ▶ Applying sort or hash method to distribute sessions into matching buckets
- ▶ Connect sessions in a same bucket

# Instance: User Tracking System




Figure: Match Features

# Instance: User Tracking System

## Claim

This optimization (using sort method) reduced time cost.

At first, the total elements to match upon each other will increase to  $O(N \times 10)$ . To sort all elements cost  $O(\log N \times 10)$ . Connect elements in a bucket cost  $O(N \times 10)$  in total. The final cost will be

$$O(\log N \times 10) \simeq O(\log N) \preceq O(N^2)$$

Thank you!